

1/6/25: Introduction

- need to install:
 - docker desktop
 - anaconda or miniconda python
 - datagrip or dbeaver
 - VS code for python development (3.10 or higher)
 - ability to interact with git and github
- should review
 - shell cmds
 - docker and docker compose
 - whats a container, whats an image, volume, mapping
- make an aws account

1/8/25: Foundations

- baseline efficiency is linear search
 - start at beginning of list and progress from the first element to the last element until you find what youre looking for or get to the last element
- record: a collection of values
 - ex. a row in a table
- collection: set of records of the same entity type
 - ex. a table
- search key: a value for an attribute from the entity type
 - ex. >1
- need $n * x$ bytes of memory
 - each record takes up x bytes
 - n is the number of records
- contiguously allocated list (array)
 - all $n * x$ bytes are allocated as a single chunk of memory
- linked list
 - each record needs x bytes + additional space for 1 or 2 memory addresses
 - takes up a bit more space
 - need some way to know if you are at the front of the list
- arrays vs linked list pros and cons
 - arrays are faster for random access but slow for inserting anywhere but the end
 - linked lists are faster for inserting anywhere in the list, but slower for random access
- binary search: cut in half each time and figure out which half target is in
 - has to be sorted
 - input: array of values in sorted order, target value
 - output: the location (index) of where target is located or some value indicating target was not found
 - usually recursive

- can be dangerous if you have a huge dataset
- time complexity
 - linear search
 - best case: $O(1)$
 - worst case: $O(n)$
 - binary search for an array (contiguously allocated list)
 - best case: $O(1)$
 - worst case: $O(\log_2 n)$
 - super inefficient on a linked list
- database searching
 - problem: can only efficiently search a table by the primary key (since it is the only sorted column)
 - we need an external data structure to support faster searching than a linear scan
 - solution: database index
- binary search tree
 - has fast insert like a linked list and fast search like an array
 - every node in the left subtree is less than its parent and every node in the right subtree is greater than its parent

1/9/25

•

1/16/25

- need to recursively descend to find all the files??
- OOP in python
 - constructor
 - add_value function appends value to a list
 - consider duplicates using a set object
 - abstract method: if a class implements this method, then implementations of the class also need the method
 - pickle it

```
def _rotate_left(self, x: AVLNode) -> AVLNode:
```

```
    T2 = y.left
```

```
    y.left = x
```

```
    x.right = T2
```

```
    ht_x_left = (0 if not x.left else x.left.height)
```

```
    ht_x_right = (0 if not x.right else x.right.height)
```

```
    x.height = 1 + (ht_x_left if ( ht_x_left >
```

1/27/25: Moving Beyond the Relational Model

- benefits of the relational model
 - mostly standard data model and query language
 - ACID compliance
 - works well with highly structured and large amounts of data
- ACID properties
 - atomicity
 - consistency
 - isolation
 - durability

2/19/25: Intro to the Graph Data Model

- graph database
 - data model based on the graph data structure
 - composed of nodes and edges
- types of graphs
 - connected (vs unconnected)
 - for every pair of nodes, there is a path
 - weighted (vs unweighted)
 - edge has a weight property
 - directed (vs undirected)
 - edges define a start and end node
 - acyclic (vs cyclic)
 - graph contains no cycles

practical

- `pip install pymupdf`
- `ollama pull mistral:latest`